

Contents

1	README	2
2	Makefile	4
3	uthreads.cpp	5

1 README

```
1  itamakatz, liavst2
2  Itamar Katz (555792977) , Liav Steinberg (203630090)
3  EX: 2
4
5  =====
6  =                FILES                =
7  =====
8
9  - README - this file.
10 - makefile
11 - uthreads.cpp
12
13 =====
14 =                ANSWERS                =
15 =====
16
17
18 Part 2: Theoretical Questions:
19 =====
20
21 1. Kernel level threads, unlike processes, share most of its context
22    with the rest of the threads. Due to that, threads are
23    fundamentally designed to work together with other threads since
24    all threads have complete access to the processes heap, code,
25    global variables etc. In other words, unlike processes, threads
26    are not independent of one another.
27    A tab in a browser could be implemented with a kernel level thread
28    but it seems pointless since there is no reason to assume tabs are
29    dependent one of the other and its not logical to give each tab
30    the ability to access and change info of other threads.
31    Besides that, a process tend to have more memory to use since by
32    implementing threads, memory is needed for the PCB.
33
34    Kernel level threads and processes have in common that they are
35    not blocked when a single process or thread get blocked (unlike
36    user level threads) and that they can take advantage of
37    multiprocessing.
38
39    Kernel level threads minimize context switching in comparison to
40    processes and thus lowers the general overhead. In addition, since
41    they share memory, kernel level threads can communication
42    considerably fast and efficiently.
43
44    On the other hand, processes tend to manage memory more
45    efficiently since there is no need of a PCB for each internal
46    thread. And as stated before, even though processes can cope with
47    each other, unlike threads they do not have to share memory and
48    information, making the process potentially autonomous and
49    independent. The latter is important since there is no impact on
50    other processes in case of failure, in contrast to threads.
51
52
53 2. We chose to send the kill signal to "pluma" (with the pid 27469).
54    The kill signal is a tool to communicate with other processes
55    using a signal. by typing "kill 27469", which by default sends the
56    SIGTERM signal to the process, we send the kill signal which by
57    definition cannot be handled by a process which then is handled by
58    the OS terminating the process.
59
```

60
61 3. Real time, is the physical time as we know it. Virtual time is
62 defined by the ability to manipulate the fixed beginning of time
63 as we wish. It can be stopped, reset, extended and done with
64 whatever we want, since it does not represent the physical world.
65 Real time can be used for example to measure the time a function
66 is executed (like we did in ex1) and virtual time can be used to
67 schedule the run of user-level threads (like we are doing in this
68 exercise).
69
70
71 4a. When using fork, a child process is spawned. The child process is
72 an exact duplicate of its parent process meaning all the data in
73 the global variables, stack and heap is identical at the spawning
74 point, but it is physically not the same memory, stored in
75 different locations, thus not shared, so from that point on, none
76 of the above is common between the processes.
77
78
79 4b. A pipe is an implementation of a communication channel for
80 processes. It contains an array with only 2 slots where one slot
81 is exclusively for writing data and the other is exclusively for
82 reading data. The array is connected to the FDs Table in a way
83 that two processes can communicate.
84 since processes are mostly autonomous and they do not share
85 memory, the pipe is needed for complex communication between
86 processes other than just signal communication like kill.

2 Makefile

```
1  CFLAGS = -g -Wall -std=c++11
2  TAR_NAME = ex2.tar
3  SOURCES = uthreads.cpp
4  HEADERS = uthreads.h
5  OSMLIB = libuthreads.a
6  OBJS = $(SOURCES:.cpp=.o)
7  EXTRA_FILES = README Makefile
8  TAR_FILES = $(SOURCES) $(EXTRA_FILES)
9
10 .DEFAULT_GOAL = $(OSMLIB)
11
12
13 all: $(OSMLIB) tar
14
15
16 $(OBJS): $(SOURCES) $(HEADERS)
17     $(CXX) -c $(CFLAGS) $<
18
19
20 $(OSMLIB): $(OBJS)
21     ar rcs $@ $^
22
23 tar:
24     tar -cvf $(TAR_NAME) $(TAR_FILES)
25
26
27 clean:
28     rm -f $(TAR_NAME) $(OBJS) $(OSMLIB)
29
30
31 .PHONY:
32     all tar clean
```

3 uthreads.cpp

```
1  using namespace std;
2
3  #include <cstdlib>
4  #include "uthreads.h"
5  #include <signal.h>
6  #include <sys/time.h>
7  #include <list>
8  #include <map>
9  #include <vector>
10 #include <queue>
11 #include <setjmp.h>
12 #include <iostream>
13 #include <memory>
14
15 #define FAILURE -1
16 #define SUCCESS 0
17 #define SYSTEM_FAILURE 1
18
19 typedef enum {READY = 1, SLEEP = 2, RUNNING = 3, BLOCKED = 4} State;
20
21 #ifdef __x86_64__
22 /* code for 64 bit Intel arch */
23
24 typedef unsigned long address_t;
25 #define JB_SP 6
26 #define JB_PC 7
27
28 /* A translation is required when using an address of a variable.
29    Use this as a black box in your code. */
30 address_t translate_address(address_t addr)
31 {
32     address_t ret;
33     asm volatile("xor    %%fs:0x30,%0\n"
34                 "rol     $0x11,%0\n"
35                 : "=g" (ret)
36                 : "0" (addr));
37     return ret;
38 }
39
40 #else
41 /* code for 32 bit Intel arch */
42
43 typedef unsigned int address_t;
44 #define JB_SP 4
45 #define JB_PC 5
46
47 /* A translation is required when using an address of a variable.
48    Use this as a black box in your code. */
49 address_t translate_address(address_t addr)
50 {
51     address_t ret;
52     asm volatile("xor    %%gs:0x18,%0\n"
53                 "rol     $0x9,%0\n"
54                 : "=g" (ret)
55                 : "0" (addr));
56     return ret;
57 }
58
59 #endif
```

```

60
61  class Thread;
62
63  // segment for the timer signal
64  struct sigaction g_sa;
65
66  // timer for the slot
67  struct itimerval g_timer;
68
69  int g_current_running_thread;
70
71  // total quanta run in the current process
72  int g_total_quantums = 0;
73
74
75  // env buffer
76  vector<sigjmp_buf> g_thread_env(MAX_THREAD_NUM);
77
78  // stack for each thread
79  vector<char[STACK_SIZE]> g_stack(MAX_THREAD_NUM);
80
81  // for blocking and unblocking the timer signal
82  sigset_t g_mask_set;
83
84  // (tid:thread) map
85  map<int, unique_ptr<Thread> > g_thread_map;
86
87  // queue of unused tids. implemented as a min_heap
88  priority_queue<int, vector<int>, greater<int> >g_unused_indexes;
89
90  // a list of all the thread tids waiting to run
91  list<int> g_ready_threads;
92
93  // sleeping map. each key is the wake up time,
94  // each value is a vector containing the all the thread
95  // tids with that wake up time.
96  map<int, vector<int> > g_sleeping_thread_map;
97
98
99
100 /**
101  * prints an error when a library function fails.
102  */
103 void terror(string msg);
104
105 /**
106  * prints an error when a system call fails, and exits.
107  */
108 void serror(string msg);
109
110 /**
111  * blocks the SIGVTALRM signal
112  */
113 inline void mask_block(){
114     if (sigprocmask(SIG_BLOCK, &g_mask_set, NULL)){
115         serror("sigprocmask failed");
116     };
117 }
118
119 /**
120  * unblocks the SIGVTALRM signal
121  */
122 inline void mask_unblock(){
123     if (sigprocmask(SIG_UNBLOCK, &g_mask_set, NULL)){
124         serror("sigprocmask failed");
125     };
126 }
127

```

```

128  /**
129   * flushes all data structures used, and exits the process.
130   */
131  void general_exit(int exit_type){
132
133      g_sleeping_thread_map.clear();
134      g_thread_map.clear();
135      g_sleeping_thread_map.clear();
136      g_ready_threads.clear();
137      g_thread_env.clear();
138      g_stack.clear();
139      mask_unblock();
140      if (sigemptyset(&g_mask_set)){
141          perror("sigemptyset failed");
142      }
143      exit(exit_type);
144  }
145
146  /**
147   * prints an error when a library function fails.
148   */
149  void terror(string msg){
150      cerr << "thread library error: " << msg << endl;
151  }
152
153  /**
154   * prints an error when a system call fails, and exits.
155   */
156  void serror(string msg){
157      cerr << "system error: " << msg << endl;
158      general_exit(SYSTEM_FAILURE);
159  }
160
161
162
163  class Thread {
164
165  private:
166      Thread() = delete;
167      State _state;
168      int _tid;
169      int _wakeup_time;
170      int _quantum_count;
171
172  public:
173
174      /**
175       * thread constructor.
176       */
177      Thread(int tid):
178          _state(READY),
179          _tid(tid),
180          _wakeup_time(0),
181          _quantum_count(0)
182      {}
183
184      /**
185       * thread destructor
186       */
187      ~Thread(){}
188
189      /**
190       * blocks a specific thread.
191       */
192      int block() {
193          if (_state == READY) {
194              _state = BLOCKED;
195              g_ready_threads.remove(_tid);

```

```

196         mask_unblock();
197         return SUCCESS;
198     }
199     else if(_state == RUNNING){
200         _state = BLOCKED;
201
202         int jump = sigsetjmp(g_thread_env[_tid], 0);
203         if(!jump){
204             check_sleepers();
205             next_run();
206             serror("siglongjmp failed");
207         }
208         return SUCCESS;
209     }
210     mask_unblock();
211     return SUCCESS;
212 }
213
214 /**
215  * resumes a specific thread from the
216  * blocked state.
217  */
218 void resume() {
219     if (_state == BLOCKED){
220         _state = READY;
221         g_ready_threads.push_back(_tid);
222     }
223     return;
224 }
225
226 /**
227  * a thread switcher. called every time the running thread
228  * is blocked / goes to sleep / the quanta ended.
229  */
230 static int next_run() {
231
232     try {
233
234         int front = g_ready_threads.front();
235         g_ready_threads.pop_front();
236
237         // next thread
238         g_current_running_thread = front;
239         if (g_thread_map.find(front) != g_thread_map.end()) {
240             g_thread_map[front]->_quantum_count++;
241             g_thread_map[front]->_state = RUNNING;
242         }
243         g_total_quantums++;
244
245
246         if (setitimer(ITIMER_VIRTUAL, &g_timer, NULL)) {
247             serror("timer failed");
248         }
249
250         mask_unblock();
251         siglongjmp(g_thread_env[front], 1);
252     }
253     catch (const out_of_range& eer) {
254         terror(eer.what());
255         general_exit(FAILURE);
256         return FAILURE;
257     }
258 }
259
260 /**
261  * puts a specific thread to sleep.
262  */
263 void sleep(int sleep_quantums) {

```



```

264     // not including the current quantum
265     if (_state == RUNNING) {
266
267         // for example went to sleep at in between slots
268         // 100-101 then wake up in slot 108 (100 + 7 + 1)
269         _wakeup_time = g_total_quantums + sleep_quantums + 1;
270         _state = SLEEP;
271         g_sleeping_thread_map[_wakeup_time].push_back(_tid);
272
273         int jump = sigsetjmp(g_thread_env[_tid], 0);
274         if(!jump){
275             check_sleepers();
276             next_run();
277             serror("siglongjmp failed");
278         }
279     }
280     mask_unblock();
281     return;
282 }
283
284 /**
285  * returns the wake up time of a specific thread.
286  */
287 int time_to_wake() {
288     if (_state == SLEEP) {
289         return (_wakeup_time - g_total_quantums);
290     }
291     return FAILURE;
292 }
293
294 /**
295  * returns the quantum count for a specific thread.
296  */
297 inline int get_quantum_count() {
298     return _quantum_count;
299 }
300
301 /**
302  * the handler for the signal.
303  */
304 static void scheduler(int){
305
306     mask_block();
307
308     check_sleepers();
309
310     g_thread_map[g_current_running_thread]->_state = READY;
311     g_ready_threads.push_back(g_current_running_thread);
312
313     int jump = sigsetjmp(g_thread_env[g_current_running_thread], 0);
314     if(!jump){
315         next_run();
316         serror("siglongjmp failed");
317     }
318 }
319
320 /**
321  * goes over the sleeping map and checks who need to wake up.
322  */
323 static void check_sleepers(){
324
325     if (g_sleeping_thread_map.find(g_total_quantums+1) ==
326         g_sleeping_thread_map.end()){
327         return;
328     }
329
330     for(vector<int>::iterator it =
331         g_sleeping_thread_map[g_total_quantums+1].begin(); it !=

```

```

332         g_sleeping_thread_map[g_total_quantums+1].end(); it++){
333
334         g_thread_map[*it]->_state = READY;
335         g_ready_threads.push_back(*it);
336     }
337     g_sleeping_thread_map.erase(g_total_quantums+1);
338 }
339
340 };
341
342
343
344 // ===== //
345
346
347
348 /*
349  * Description: This function initializes the Thread library.
350  * You may assume that this function is called before any other Thread library
351  * function, and that it is called exactly once. The input to the function is
352  * the length of a quantum in micro-seconds. It is an error to call this
353  * function with non-positive quantum_usecs.
354  * Return value: On success, return 0. On failure, return -1.
355  */
356 int uthread_init(int quantum_usecs){
357
358     if(quantum_usecs > 0){
359
360         // init the unused indexes queue
361         for (int i = 0; i < MAX_THREAD_NUM; ++i) {
362             g_unused_indexes.push(i);
363         }
364
365         int tid = g_unused_indexes.top();
366         g_unused_indexes.pop();
367
368         g_thread_map[tid] = unique_ptr<Thread> (new Thread(tid));
369         g_ready_threads.push_back(tid);
370
371         // setting function where SIGVTALRM signal goes
372         g_sa.sa_handler = &Thread::scheduler;
373         g_sa.sa_flags = 0;
374
375         if (sigaction(SIGVTALRM, &g_sa, NULL)) {
376             error("sigaction failed");
377         }
378
379         // initiate the time. set to one quantum
380         g_timer.it_value.tv_sec = 0;
381         g_timer.it_value.tv_usec = quantum_usecs;
382
383         g_timer.it_interval.tv_sec = 0;
384         g_timer.it_interval.tv_usec = quantum_usecs;
385
386         if (sigemptyset(&g_mask_set)){
387             error("sigemptyset failed");
388         }
389         if (sigaddset(&g_mask_set, SIGVTALRM)){
390             error("sigaddset failed");
391         }
392         if (sigprocmask(SIG_SETMASK, &g_mask_set, NULL)) {
393             error("sigprocmask failed");
394         }
395
396         int jump = sigsetjmp(g_thread_env[tid], 0);
397         if(!jump){
398             Thread::next_run();
399             error("siglongjmp failed");

```

```

400     }
401     return SUCCESS;
402 }
403
404     terror("non-positive quantum time");
405     return FAILURE;
406 }
407
408 /*
409  * Description: This function creates a new Thread, whose entry point is the
410  * function f with the signature void f(void). The Thread is added to the end
411  * of the READY threads list. The uthread_spawn function should fail if it
412  * would cause the number of concurrent threads to exceed the limit
413  * (MAX_THREAD_NUM). Each Thread should be allocated with a stack of size
414  * STACK_SIZE bytes.
415  * Return value: On success, return the ID of the created Thread.
416  * On failure, return -1.
417 */
418 int uthread_spawn(void (*f)(void)){
419
420     mask_block();
421
422     if (g_unused_indexes.empty()){
423         // no more threads can be made (by the given limit)
424         terror("thread limit exceeded");
425         mask_unblock();
426         return FAILURE;
427     }
428
429     int tid = g_unused_indexes.top();
430     g_unused_indexes.pop();
431
432     g_thread_map[tid] = unique_ptr<Thread> (new Thread(tid));
433
434     sigsetjmp(g_thread_env[tid], 0);
435     (g_thread_env[tid]->__jmpbuf)[JB_SP] = translate_address
436         ((address_t)(g_stack[tid]) + STACK_SIZE - sizeof(address_t));
437     (g_thread_env[tid]->__jmpbuf)[JB_PC] = translate_address((address_t)f);
438
439     g_ready_threads.push_back(tid); // now it is ready to run
440
441     mask_unblock();
442     return tid;
443 }
444
445 /*
446  * Description: This function terminates the Thread with ID tid and deletes
447  * it from all relevant control structures. All the resources allocated by
448  * the library for this Thread should be released. If no Thread with ID tid
449  * exists it is considered as an error. Terminating the main Thread
450  * (tid == 0) will result in the termination of the entire process using
451  * exit(0) [after releasing the assigned library memory].
452  * Return value: The function returns 0 if the Thread was successfully
453  * terminated and -1 otherwise. If a Thread terminates itself or the main
454  * Thread is terminated, the function does not return.
455 */
456 int uthread_terminate(int tid){
457
458     mask_block();
459
460     if(tid == 0){
461         general_exit(SUCCESS);
462     }
463
464     if (g_thread_map.find(tid) == g_thread_map.end()){
465         terror("terminate: thread not found");
466         mask_unblock();
467         return FAILURE;

```

```

468     }
469
470
471     g_thread_map.erase(tid);
472     g_unused_indexes.push(tid);
473     if (tid == g_current_running_thread){ // the thread terminated itself
474         Thread::check_sleepers();
475         Thread::next_run();
476         serror("siglongjmp failed");
477     }
478
479     g_ready_threads.remove(tid);
480     mask_unblock();
481     return SUCCESS;
482
483 }
484
485 /*
486 * Description: This function blocks the Thread with ID tid. The Thread may
487 * be resumed later using uthread_resume. If no Thread with ID tid exists it
488 * is considered as an error. In addition, it is an error to try blocking the
489 * main Thread (tid == 0). If a Thread blocks itself, a scheduling decision
490 * should be made. Blocking a Thread in BLOCKED or SLEEPING states has no
491 * effect and is not considered as an error.
492 * Return value: On success, return 0. On failure, return -1.
493 */
494 int uthread_block(int tid){
495
496     mask_block();
497
498     if(tid == 0){
499         terror("main thread cannot be blocked ");
500         mask_unblock();
501         return FAILURE;
502     }
503
504     if (g_thread_map.find(tid) == g_thread_map.end()) {
505         terror("block: thread not found");
506         mask_unblock();
507         return FAILURE;
508     }
509
510     return g_thread_map[tid]->block();
511
512 }
513
514 /*
515 * Description: This function resumes a blocked Thread with ID tid and moves
516 * it to the READY state. Resuming a Thread in the RUNNING, READY or SLEEPING
517 * state has no effect and is not considered as an error. If no Thread with
518 * ID tid exists it is considered as an error.
519 * Return value: On success, return 0. On failure, return -1.
520 */
521 int uthread_resume(int tid){
522
523     mask_block();
524
525     if (g_thread_map.find(tid) == g_thread_map.end()) {
526         terror("resume: thread not found");
527         mask_unblock();
528         return FAILURE;
529     }
530
531     g_thread_map[tid]->resume();
532     mask_unblock();
533     return SUCCESS;
534
535 }

```

```

536
537 /*
538  * Description: This function puts the RUNNING Thread to sleep for a period
539  * of num_quantums (not including the current quantum) after which it is moved
540  * to the READY state. num_quantums must be a positive number. It is an error
541  * to try to put the main Thread (tid==0) to sleep. Immediately after a Thread
542  * transitions to the SLEEPING state a scheduling decision should be made.
543  * Return value: On success, return 0. On failure, return -1.
544  */
545 int uthread_sleep(int num_quantums) {
546     mask_block();
547
548     if (num_quantums <= 0){
549         terror("cannot sleep for negative time");
550         mask_unblock();
551         return FAILURE;
552     }
553
554     if (g_current_running_thread == 0) {
555         terror("main thread cannot sleep");
556         mask_unblock();
557         return FAILURE;
558     }
559     else{
560         g_thread_map[g_current_running_thread]->sleep(num_quantums);
561         return SUCCESS;
562     }
563
564     mask_unblock();
565     return SUCCESS;
566 }
567
568 /*
569  * Description: This function returns the number of quantums until the Thread
570  * with id tid wakes up including the current quantum. If no Thread with ID
571  * tid exists it is considered as an error. If the Thread is not sleeping,
572  * the function should return 0.
573  * Return value: Number of quantums (including current quantum) until wakeup.
574  */
575 int uthread_get_time_until_wakeup(int tid){
576     mask_block();
577
578     if (g_thread_map.find(tid) == g_thread_map.end()){
579         terror("wake up: thread not found");
580         mask_unblock();
581         return FAILURE;
582     }
583
584     int return_val = g_thread_map[tid]->time_to_wake();
585     if(return_val < 0) {
586         mask_unblock();
587         return SUCCESS;
588     }
589     mask_unblock();
590     return return_val;
591 }
592
593 /*
594  * Description: This function returns the Thread ID of the calling Thread.
595  * Return value: The ID of the calling Thread.
596  */
597 int uthread_get_tid(){
598     return g_current_running_thread;
599 }
600
601
602
603

```

```

604  /*
605  * Description: This function returns the total number of quantums that were
606  * started since the library was initialized, including the current quantum.
607  * Right after the call to uthread_init, the value should be 1.
608  * Each time a new quantum starts, regardless of the reason, this number
609  * should be increased by 1.
610  * Return value: The total number of quantums.
611  */
612  int uthread_get_total_quantums(){
613      return g_total_quantums;
614  }
615
616  /*
617  * Description: This function returns the number of quantums the Thread with
618  * ID tid was in RUNNING state. On the first time a Thread runs, the function
619  * should return 1. Every additional quantum that the Thread starts should
620  * increase this value by 1 (so if the Thread with ID tid is in RUNNING state
621  * when this function is called, include also the current quantum). If no
622  * Thread with ID tid exists it is considered as an error.
623  * Return value: On success, return the number of quantums of the Thread with
624  * ID tid. On failure, return -1.
625  */
626  int uthread_get_quantums(int tid){
627
628      mask_block();
629
630      if (g_thread_map.find(tid) == g_thread_map.end()){
631          terror("get quantum: thread not found");
632          mask_unblock();
633          return FAILURE;
634      }
635
636      mask_unblock();
637      return g_thread_map[tid]->get_quantum_count();
638
639  }

```